SYSTEM WALKTHROUGH

DocQuorum

From document to structured JSON

The task: extract the fields and repeated records requested by a schema from a PDF or image, while preserving what the document actually prints.

DocQuorum prepares document evidence, chooses one document route, and produces one or more separate candidate JSONs. It checks each candidate independently and, when the route uses quorum, selects the candidate with the strongest agreement with the others.

1 Define the extraction task

The input is a document D, a schema S (or a JSON template), and optionally a task description T. The intended output is a JSON object Y whose fields, types, and repeated records follow S and whose values are supported by D:

$$(D, S, T) \longrightarrow Y.$$

The schema supplies both structure and instructions. For example, it can request a purchase-order number, a supplier name, and every line item with its quantity and unit price. Descriptions can specify grouping, units, missing values, or whether a page number means the PDF index or the printed folio.

	Example input	Intended output
Document	Purchase order with a header and item table	Values grounded in those pages
Schema	Order number; supplier; list of items	One object with those keys and an item array
Optional task	“Extract every ordered item.”	Task text becomes the schema’s top-level description

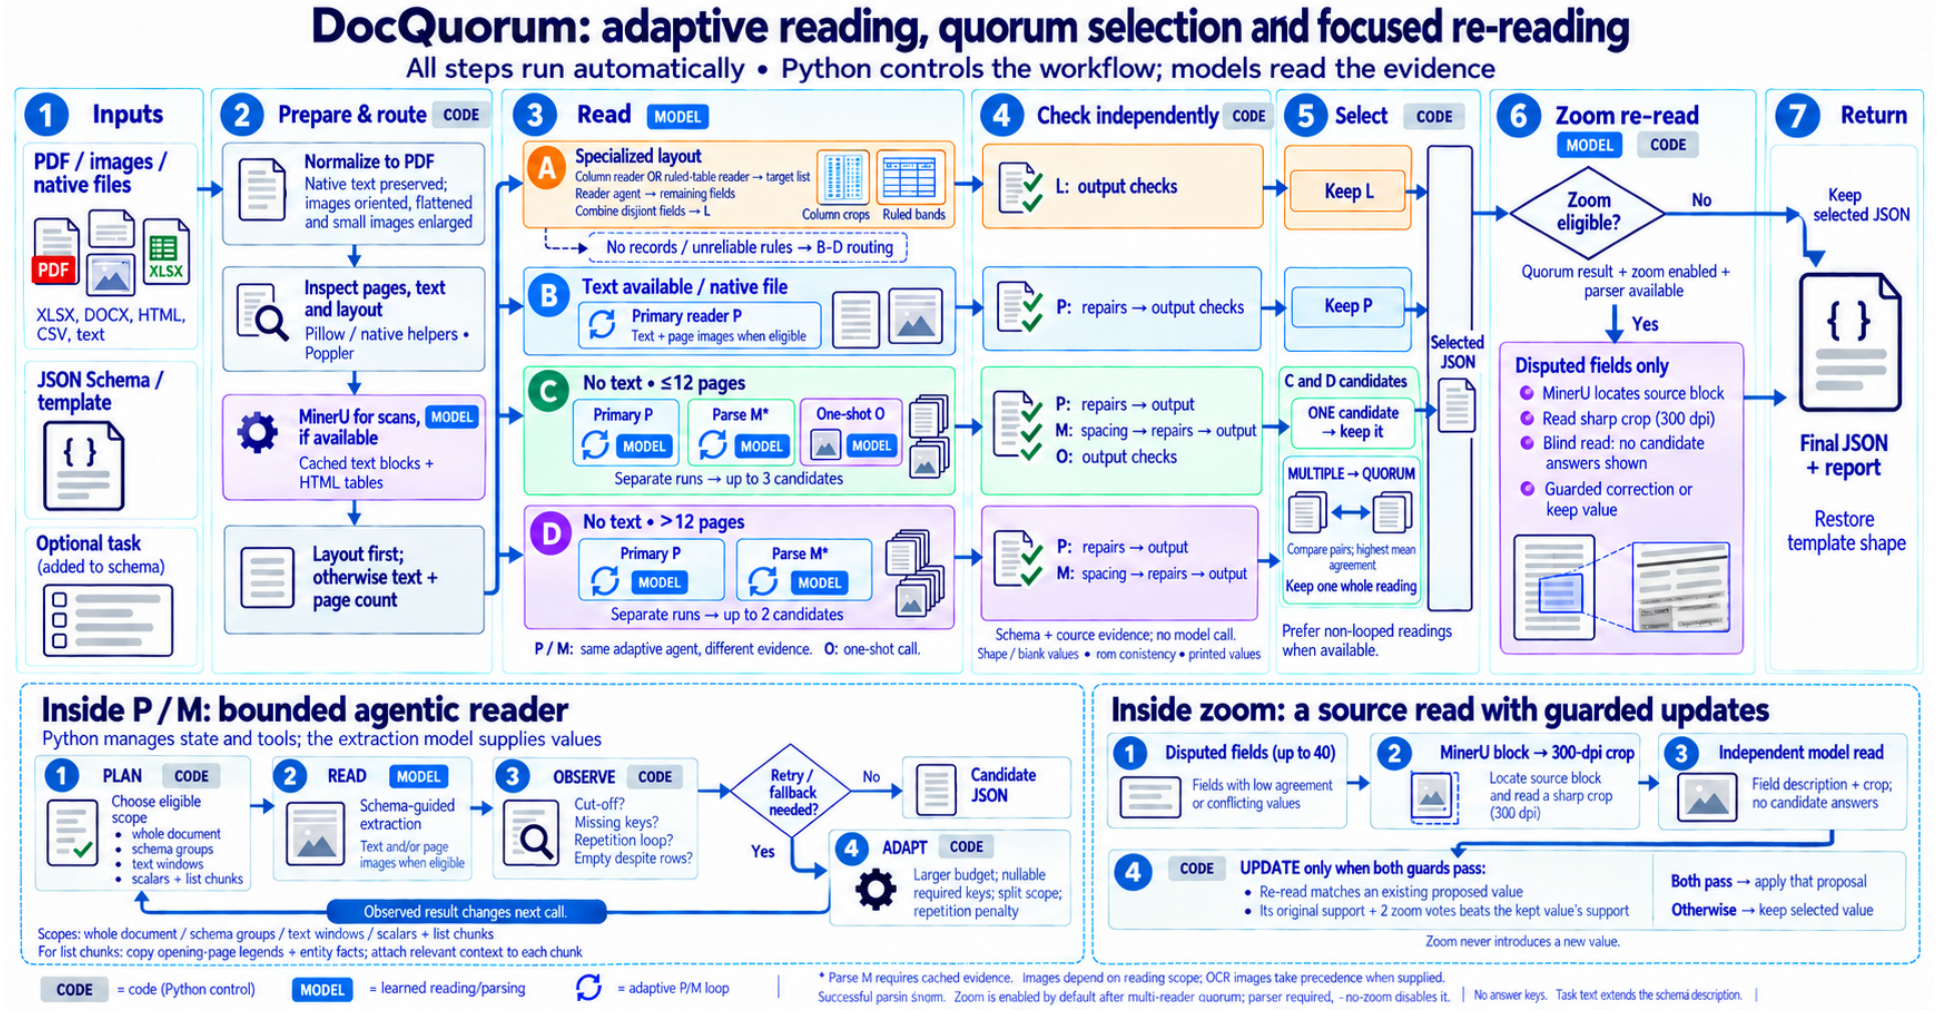
What a “reading” means

A reading produces one candidate JSON, using one model call or many smaller calls. MinerU produces source evidence: page text and HTML tables. The parse-based reader maps that evidence to the schema. Readers can share source evidence, but never receive one another’s candidate JSONs.

How to read this guide

Start with the route overview on page 2, then the evidence and candidate flows on page 3. Pages 6–7 show complete schemas and answers. Page 11 follows a purchase order through three readings; page 12 compares document routes. The results section starts on page 14.

Detailed examples originate at `b4ea9a1`; the main overview follows v1.3. The current specification is `docs/DESIGN.md`. Benchmark differences are explicit on page 13; the results tables follow on page 14.

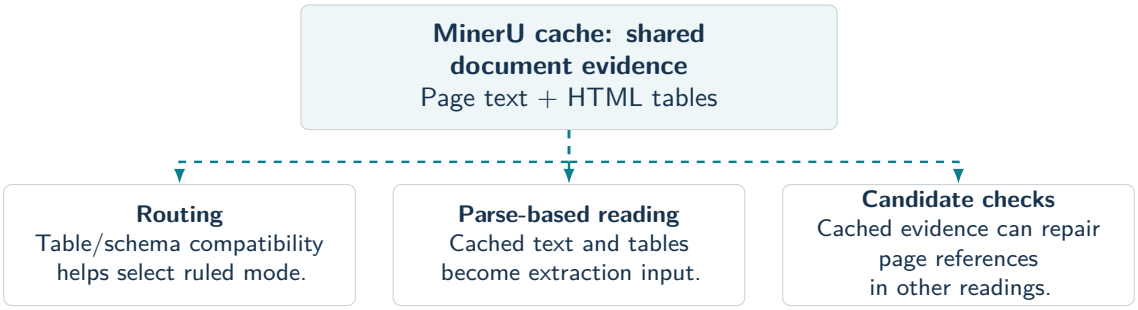


Reader components and tools. Python controls execution and the bounded P/M feedback loop. The figure shows successful paths; runtime exceptions are omitted. The v1.3 overview includes guarded zoom re-reading after quorum; other worked examples below retain their original run provenance. Source roles and eligibility details: `docs/AGENTS_AND_TOOLS.md`.

SHARED EVIDENCE / SEPARATE ANSWERS

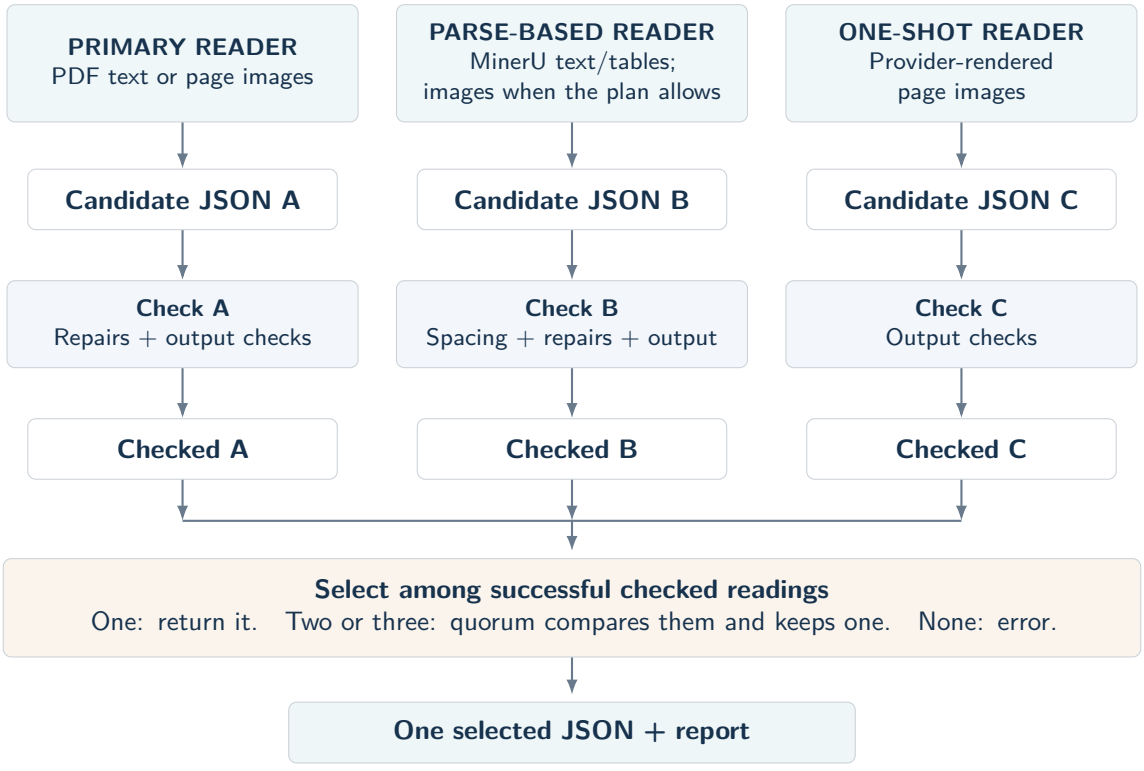
Where MinerU helps

MinerU's cache can help routing, the parse-based reader, and checks on other readings. The primary and one-shot extraction calls do not receive MinerU text. Primary and parse-based readings are separate executions of the same general reader code with different evidence.



SHORT SCAN EXAMPLE / NO LAYOUT SPECIALIST

Same document and schema; up to three separate readings. Each sees the evidence shown below.



Dashed arrows: MinerU source evidence. Solid arrows: candidate flow.

No reader-to-reader answer flow. Checks inspect source evidence, not a peer's JSON. The API attempts readers in sequence; these lanes do not imply concurrent execution.

Example of indirect help: a primary candidate says `source_page: 1`. If the schema asks for printed numbering and the cached parse of PDF page 1 ends with "7", its page-index check can change the value to 7. It uses the source cache, not the parse reader's candidate JSON.

STEPS 1–2 / PREPARE AND ROUTE

Turn the document into usable evidence

Step 1. Normalize and inspect

Input: PDF or image, schema/template, optional task.

Operation: copy a PDF into a temporary workspace, or convert an image to an RGB PDF at 150 dpi. Convert a template to a schema. If supplied, the task replaces the top-level schema description. Count pages and inspect the PDF text.

The routing test is `len(plain_text.strip()) >= 50 * max(1, n_pages)`. It measures aggregate text availability, not OCR accuracy. A scanned PDF with enough embedded OCR text takes the text-present route. Mixed documents are classified as a whole.

For a document below this threshold, the API probes the parser endpoint. If available, MinerU processes the pages and caches each page as `NNNN.txt`. Text blocks remain text; tables retain HTML cell structure. Existing cached files are reused. The document cache key is derived from its bytes.

Output: normalized PDF, page count, text flag, text-layer strings, optional cached parse. Page images are rendered when a reader requests them.

Step 2. Select a reader route

Priority	Condition	Decision
1. Column flow	At least two pages; at least three nearly equal-width columns on at least half of sampled pages	Column reader
2. Ruled scan	No sufficient text layer; cached tables map to schema fields with conversion ratio ≥ 0.9	Ruled reader
3. General	Neither specialized condition holds	General reading route

The column test samples up to four pages and requires column-width coefficient of variation below 0.12. Ruled mode is intended for scanned tables with printed lines between records. Its router first tests the cached table's compatibility with the schema; the reader later detects the printed lines in page images.

What "conversion ratio" counts

Conversion means moving parsed table cells into records with the requested field names. For example, a table row `item: Pen, quantity: 5` becomes a JSON record with `item` and `quantity` fields. The router computes:

$$\text{conversion ratio} = \frac{\text{converted records}}{\text{considered parsed table-body rows}}.$$

If 95 records are converted from 100 considered rows, the ratio is 0.95 and passes the 0.9 threshold. Pages marked bad by the converter are excluded from this calculation. The ratio measures structural compatibility; it does not measure value accuracy or coverage of every document page.

The cache remains source evidence. The ruled reader extracts its final records from image bands, with a detected row count and the table header attached to each band. **Data boundary:** answer keys do not enter preparation, routing, reading, checks, or candidate selection. Benchmark scoring is a separate operation after extraction.

STEP 3 / GENERAL READER

Plan, read, inspect, and reduce the scope

Input: schema and PDF, with either text-layer pages or cached parse pages. The primary reader uses page images when text is insufficient. Parse mode uses cached text only if its page count covers at least 90% of the extracted page-text count; otherwise it falls back to the text layer.

Output: one candidate JSON and call/plan reports. Both primary and parse-based readings use these fallback strategies: try eligible plans in order and stop at an accepted result.

Plan	When attempted	What it does
1. Whole document	Input estimate $\leq$ 100,000 tokens; output estimate $\leq$ 140,000; image-only input has $\leq$ 12 pages	**Keep everything together.** Full schema + available document evidence in one call.
2. Schema groups	Whole plan did not return; schema exceeds 60,000 compact characters; text exists or $n \leq 12$	**Split requested fields.** Read smaller schema groups; use page windows if needed.
3. Text windows	Earlier plans did not return; text exists; no top-level arrays	**Split source text.** Ask for the same fields in successive windows; merge filled values.
4. Scalars + lists	Earlier plans did not return	**Read scalars and lists.** Read single-value fields, then each list in page chunks.

An invoice's number, date and supplier are scalars; its item rows form a list. The 140,000-token output estimate is a planning cutoff, not a response budget: actual whole-document replies are capped at 32,000 tokens in this baseline.

What each model call sees

Whole calls attach up to 12 page images; parse-based scans of 13–24 pages can pair pages. Longer whole calls use text only. Partitioned long scans can use eight images per window; ordinary text windows target 60,000 tokens. OCR is labeled as machine-read alongside images.

Text chunks contain text only; image chunks attach up to two pages. Chunk size is one or two pages, based on row density and field count against a 9,000-token output target. Fields named with "page" force one-page chunks. Earlier chunk predictions are not fed into later calls; rows are collected in page order.

What causes a retry or fallback

- A whole reply ending with `finish_reason=length` may be retried at 32,000 output tokens. Continued truncation triggers the next plan.
- A guided whole reply missing schema keys is retried with all keys required; a usable retry replaces it.
- Truncated *text* chunks split recursively, up to depth four; keep the halves' result only if it has more rows. Image chunks are not recursively split by this logic.
- A text chunk with at least five estimated rows but no extracted records gets one extra attempt with a row-count reminder.

Schema/window merging fills empty fields, keeps the first nonempty scalar and the longer list; window lists are not concatenated. Scalar extraction uses the first two pages in image mode; for text over 240,000 characters, it uses the first three and last two pages.

SCHEMA + ANSWER / PURCHASE ORDER

Single fields and repeated records

This synthetic source is the purchase order used in the walkthrough on page 11. The schema and answer below are complete JSON, with no omitted fields or rows.

Source header: PO-1042 Supplier: North Paper

Ordered item	Quantity	Unit price
A4 paper	10	4.50
Envelopes	2	6.00

INPUT / JSON SCHEMA

```
{
  "$schema": "http://json-schema.org/draft-07/schema
    #",
  "type": "object",
  "description": "Extract the order header and every
     ordered item.",
  "properties": {
    "order_number": {"type": "string", "description":
     "The purchase-order number as printed."},
    "supplier": {"type": "string", "description": "
    The supplier name as printed."},
    "items": {
      "type": "array",
      "description": "Every ordered item, in
    document order.",
      "items": {
        "type": "object",
        "properties": {
          "description": {"type": "string"},
          "quantity": {"type": "number"},
          "unit_price": {"type": "number"}
        },
        "required": ["description", "quantity", "
    unit_price"]
      }
    }
  },
  "required": ["order_number", "supplier", "items"]
}
```

EXPECTED ANSWER / JSON

```
{
  "order_number": "PO-1042",
  "supplier": "North Paper",
  "items": [
    {"description": "A4 paper", "
    quantity": 10, "unit_price": 4.5},
    {"description": "Envelopes", "
    quantity": 2, "unit_price": 6.0}
  ]
}
```

Scalars

`order_number` and `supplier` each hold one value.

List

`items` holds two records. Each record has its own description, quantity and unit price.

If the scalars + lists plan runs

Read the scalar fields separately; read item rows in page chunks; combine them into the one candidate shown above. Checks happen afterward.

Schema versus answer: `properties` names the requested fields; `type` sets their data types; `items` defines the structure of each array entry; `required` lists keys that must be present. The answer contains values, not these schema declarations.

This small source may fit the whole-document plan. The example demonstrates the scalar/list split, not a measured execution of the fallback plan.

A real electricity-bill task

DocuBench task `PSU5pciM`: a two-page CPS Energy sample bill. Both JSONs below retain the complete published content, with whitespace reformatted. The answer is the benchmark reference, not a new DocQuorum prediction.

INPUT / FULL PUBLISHED SCHEMA

```
{
  "$schema": "http://json-schema.org/draft-07/schema#",
  "type": "object",
  "description": "residential electric utility bill with a charges breakdown and a
    meter reading",
  "properties": {
    "billingPeriodStart": {"type": "string", "description": "the start date of the
    billing period in YYYY-MM-DD."},
    "billingPeriodEnd": {"type": "string", "description": "the end date of the billing
    period in YYYY-MM-DD."},
    "previousMeterRead": {"type": "number", "description": "the previous meter reading
    ."},
    "currentMeterRead": {"type": "number", "description": "the current meter reading
    ."},
    "consumptionKwh": {"type": "number", "description": "the metered electricity
    consumption in kWh for the period."},
    "charges": {
      "type": "array",
      "description": "the electric charge line items, each a charge name and its amount
    . exclude the total/summary row.",
      "items": {
        "type": "object",
        "properties": {
          "description": {"type": "string", "description": "the charge name as printed
    on the line."},
          "amount": {"type": "number", "description": "the charge amount."}
        }
      }
    },
    "totalElectricBill": {"type": "number", "description": "the total electric bill
    amount, equal to the sum of the charge lines."}
  }
}
```

EXPECTED ANSWER / ALL FOUR CHARGES

```
{
  "billingPeriodStart": "2021-09-18",
  "billingPeriodEnd": "2021-10-18",
  "previousMeterRead": 47911.0,
  "currentMeterRead": 48374.0,
  "consumptionKwh": 463.0,
  "charges": [
    {"description": "Service Availability Charge", "amount":
    8.75},
    {"description": "Energy Charge 463 kWh x $0.0691", "amount":
     31.99},
    {"description": "Fuel Adjustment 463 kWh x $0.02568", "
    amount": 11.89},
    {"description": "Regulatory Adj 463 kWh x $0.01236", "amount
    ": 5.72}
  ],
  "totalElectricBill": 58.35
}
```

Source facts, reconstructed from the bill:
Billing period: Sep 18–Oct 18, 2021. Meter readings: previous 47,911; current 48,374; consumption 463 kWh. Electric charges: $8.75, $31.99, $11.89 and $5.72. Total: $58.35.

Six scalars: the two dates, two meter readings, consumption and total. **One list:** `charges`, with a description and amount for each row.

The total is kept outside `charges`; it is not a fifth charge. Printed dates become ISO date strings, and amounts become numbers without dollar signs.

Schema and reference answer: DocuPipe / DocuBench, revision `43a3f3b`; CC BY 4.0. Whitespace changes only. Source document: CPS Energy redacted sample.

STEP 3 / ALTERNATIVE READERS

Read the geometry that defines a record

Column reader: preserve the reading order

Input: page images from a document routed as column flow, and the schema's main array.

Operation:

1. Render each page at 300 dpi and crop its detected columns.
2. Read each crop from top to bottom with the item schema. The decoding schema requires the item fields and permits a null lead field.
3. Collect records in page order, then column order. Assign the physical PDF page index to `source_page`.
4. If a record has an empty lead field and a preceding record exists, treat it as a continuation: fill empty fields in the previous record and append string content.

Output: one candidate object containing the main array. The lead field is the first required string field whose name does not contain "page", with a fallback to the first property. A mistaken empty lead field can therefore cause an incorrect join.

Ruled reader: use printed lines to count records

Input: a routed scanned table and the schema's main array.

Operation:

1. Render the page at 200 dpi, deskew it, and detect horizontal rules.
2. Identify a header and row bands of up to 12 rows.
3. Stack the header above each band so every call sees the column labels.
4. Decode an array with `minItems = maxItems = n`, where n is the detected row count for that band.
5. Attach the physical page index and concatenate records in page order.

Output: one main-array candidate, then the layout check chain. A page with no detected row bands yields no records from this reader. The row constraint relies on correct rule detection; it does not establish that the values are correct.

One-shot reader: an additional candidate for short scans

The general API attempts the ExtractBench `VLLMExtractProvider` reader for documents without sufficient text and with at most 12 pages. It uses that provider's prompt and rendering configuration (150 dpi) with the configured extraction model. Its JSON receives the one-shot check chain.

This is a separate reading strategy. It does not revise the primary reader's JSON. If this call raises an exception, the API logs the failure and continues with available candidates.

Why these readers are separate

A column crop makes small print readable and preserves column order. A ruled band provides a physical count of rows. The one-shot strategy provides another interpretation of the same short scan. Specialized layout output is returned directly after checks; it does not compete in the general quorum.

STEP 4 / CHECK EACH CANDIDATE

Repair JSON using the schema and evidence

Input: one candidate JSON and a `checks.Doc` object containing the schema and lazily loaded text layer, reading-order text, and cached parse.

Operation: run deterministic transformations in a fixed order. Every transformation receives a fresh JSON copy. There are no model calls or answer keys.

Output: one checked candidate, still separate from every other reading.

All chains can access the same document evidence. In particular, `page_index` can use the cached parse in primary/parse repairs, and `printed_pages` can use it in any chain. A primary or one-shot candidate can therefore benefit from MinerU during checks, although its extraction call did not receive the parse. Checks never receive another candidate's JSON.

Candidate source	Exact chain
Primary	Repairs → Output checks
Parse	Spacing → Repairs → Output checks
One-shot / layout	Output checks only

Spacing and repairs, in order

Module	Intended transformation; applies only when its guards match
`spacing`	Restore text-layer spacing when a parsed string differs only in spaces.
`carry_down`	Propagate a group value to later rows when the schema describes carry-down behavior.
`page_index`	Normalize `source_page`; translate physical indices to printed folios where requested.
`dedupe`	Deduplicate on a key identified by the list's uniqueness instructions.
`stacked_rows`	Join records split across multiple printed lines as described by the schema.

Output checks, in order

Module	Intended transformation; applies only when its guards match
`conform`	Fill nullable properties; repair missing containers and null/empty-list conventions.
`blank_boxes`	Set a blank checkbox to false when the field description requires it.
`template_rows`	Remove empty preprinted rows or template records.
`copied_rows`	Remove records spuriously copied into another array.
`record_counts`	Cap duplicate records using repetition found in document text.
`printed_pages`	Locate rows through distinctive words and correct printed page labels.
`printed_quotes`	Restore uniquely supported quote/apostrophe forms from the text layer.
`printed_numbers`	Recover uniquely supported printed precision across unit scales; integer-only fields stay unchanged.

These are conditional heuristics. Some cannot act without suitable text evidence. `conform` is not a full JSON Schema validator, and a checked object is not a certificate of factual correctness. For example, missing required scalar values are not universally repaired.

STEP 5 / CHOOSE A READING

Agreement selects one complete candidate

Input: checked candidates for the same document and schema.

Policy: a layout route returns its checked result; a scan longer than 12 pages returns its checked parse reading when available. Otherwise, one candidate is returned directly, several invoke the quorum, and no candidates raises an error.

The quorum calculation

For candidate c, treat each other candidate o as a temporary reference. The ExtractBench scorer computes value precision $P(c \mid o)$ and recall $R(c \mid o)$ under the schema. The selector explicitly supplies empty field rules and repeated-structure rules; it does not use the benchmark's annotated answer or custom matching rules.

$$F_2(c \mid o) = \frac{5P(c \mid o)R(c \mid o)}{4P(c \mid o) + R(c \mid o)}, \qquad A(c) = \frac{1}{m-1}\sum_{o \neq c} F_2(c \mid o).$$

If precision or recall is nonpositive, the pair contributes zero. The candidate with maximum $A(c)$ wins. Scores are rounded to nine decimals before comparison; ties go to insertion order (primary, parse, one-shot in the API).

Recall is weighted more heavily: a candidate is rewarded for retaining values found by other readers. Pairwise comparisons are directional. The selected JSON remains a complete checked candidate; fields are not voted on or stitched together.

Worked selection example

Consider a simplified array where every item is a distinct matching value. Candidate A contains `[a,b]`, B contains `[a,b,c]`, and C contains `[a,b,c]`. This illustrates the arithmetic; real structured matching uses the scorer.

Candidate	Against first peer	Against second peer	Mean F_2
A	B: $P = 1, R = 2/3$	C: $P = 1, R = 2/3$	0.714286
B	A: $P = 2/3, R = 1$	C: $P = 1, R = 1$	0.954545
C	A: $P = 2/3, R = 1$	B: $P = 1, R = 1$	0.954545

B wins the tie with C because B is earlier. Its complete array is returned. Agreement cannot reveal an error shared by all readings; a high agreement score is not a calibrated probability of correctness.

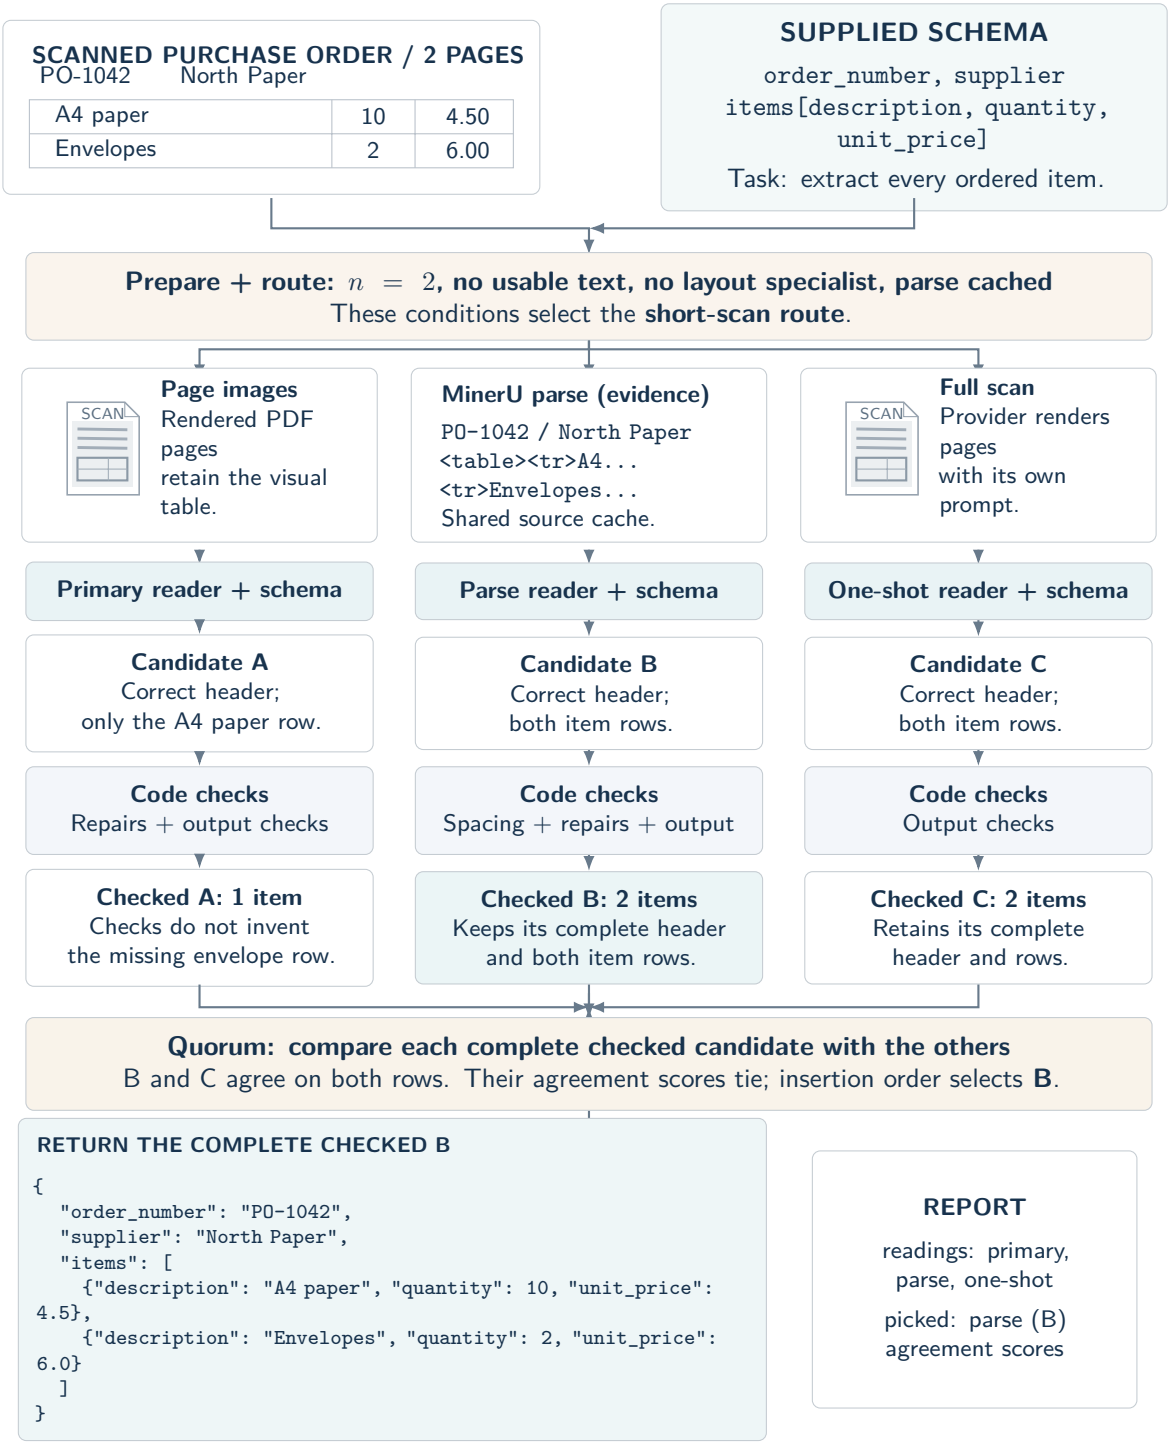
Step 6. Format and return

For a JSON Schema input, return the selected object. For a template input, `fill_template` restores the original template shape and uses empty strings for absent leaves. The Python API returns `(result, report)`. A multi-candidate report includes candidate names, the pick, and rounded agreement scores. Direct returns include a readings list; the long-scan direct return reports only "parse", even though primary was attempted first.

END-TO-END EXAMPLE

Follow the evidence into the final JSON

This is a synthetic walkthrough, not a benchmark result. Assume there is no sufficient text layer, no specialized layout route, and the parser is available. Each reader produces its own candidate; the schema is common to all three.



Illustrative values and candidate errors. Checks may consult the shared source cache, never another candidate's JSON. The API attempts readers in sequence. Quorum does not combine A's header with B's rows.

ROUTING COMPARISONS

Four documents, four different paths

Follow one pictured case at a time. The layout test comes first; text availability and page count decide the remaining routes.

DIGITAL PDF OR SCAN
ENOUGH EMBEDDED TEXT



No specialist selected.
Enough embedded text selects the text route, even if the PDF looks scanned.

Primary reader → Primary checks → Return its JSON

Page images may accompany text. No general quorum is needed for this single reading.

SHORT SCAN
PARSER UNAVAILABLE



No sufficient text; $n \leq 12$.
No parse candidate, but primary and one-shot may still succeed.

Primary + one-shot → Check each separately → Quorum if both succeed

One successful reading returns directly. No successful reading raises an error.

LONG SCAN
SUCCESSFUL CACHED PARSE



No sufficient text; e.g. $n = 20$.
Try primary, then the parse reading.
A successful parse takes priority.

Read primary, then parse → Check each separately → Return checked parse

No one-shot or quorum. Primary is the fallback if the parse reading is unavailable.

SPECIALIZED LAYOUT
SELECTED FIRST



Equal-width columns or a routed ruled scan. This decision has priority over the general text/scan routes.

Column / ruled reader → Layout checks → Return its JSON

One geometry-aware candidate. The API does not compare it with general-reader candidates.

One document route is chosen. The short-scan route can run several readings, each with a separate candidate JSON. MinerU evidence can help routing and checks as well as the parse-based reader; no reader sees another reader's answer. A high agreement score cannot detect an error shared by every candidate.

EXECUTION AND SCOPE

Connect the guide to the code

Run a document through the API

From the repository root, with dependencies and model endpoints configured as described in the repository README:

```
python -m docquorum.extract invoice.pdf \
  --schema invoice_schema.json -o invoice.json
```

```
from docquorum.extract import extract
result, report = extract("invoice.pdf", schema_dict)
```

The CLI writes the JSON to the chosen file (or stdout) and the final report to stderr. Some reader diagnostics use direct prints, so `-o` is the reliable way to capture a separate JSON artifact.

Implementation map

Source under `docquorum/`	Responsibility
`extract.py`	Preparation, route, candidate lifecycle, final formatting
`tools/pdf.py`, `tools/parser.py`	Evidence extraction, rendering, cached page parses
`agents/layout.py`	Specialized layout decision
`agents/reader.py`, `agents/whole.py`	General reader plans, retries, chunks, windows
`agents/column_reader.py`, `agents/ruled_reader.py`	Specialized visual readings
`agents/oneshot.py`	Additional short-scan provider
`checks/__init__.py`, `checks/*.py`	Exact ordered chains and transformations
`quorum.py`	Pairwise agreement and whole-candidate selection
`tools/schema.py`	Schema decomposition and template formatting

Benchmark runners have different orchestration

ExtractBench: `benchmarks/extractbench/run.sh` runs primary over the split first, then parse for scans, then one-shot/quorum for short scans. Layout results overwrite general results on medium/long splits only. The script reads cached parses; prepare them with `python -m docquorum.tools.parser SPLIT`. Unlike the API, it does not create missing parses automatically. Its quorum command may also copy parse-only long-scan results. Scoring uses answer keys only after extraction.

UNIKIE: `benchmarks/unikie/unikie_eval.py` uses a separate template-based image protocol and six readings: image, one-shot plus parse text, one-shot, parse, and two extra-model one-shots. It uses normalized flattened-value matching for its F_2 quorum. It does not execute the API's layout/reader/check pipeline.

Practical limits

Routing thresholds, row counts, and text repairs are heuristics. Parser caches can contain empty pages. Some call failures can leave partial results; the API excludes general-reader results with recorded server errors, but this does not detect every incomplete output. Specialist failures are not caught by a general-reader fallback. Inspect important extractions against the original document; neither structured decoding nor agreement proves completeness.

RESULTS

ExtractBench

370 enterprise PDFs; value F1 (%) by the benchmark's own scorer. Published systems: ExtractBench leaderboard (2 October 2026). Equal scores share a rank.

#	System	Short	Medium	Long	Overall
1	LlamaExtract Agentic Plus	96.94	94.45	90.18	95.91
1	Pulse (Effort)	96.46	95.01	93.51	95.91
3	**DocQuorum (ours)**	**96.34**	**94.11**	**90.41**	**95.43**
4	LlamaExtract Agentic	96.04	92.11	91.83	94.77
5	Codex (GPT-6 Sol Evidence)	95.99	92.42	87.25	94.57
6	Codex (GPT-5.6 Sol Evidence)	96.04	90.19	82.73	93.77
7	Codex (GPT-5.5)	95.68	91.15	78.88	93.57
8	Codex (GPT-5.5 Evidence)	95.59	88.72	87.97	93.35
9	Pulse (Non-effort)	93.17	92.22	89.07	92.70
10	LlamaExtract Cost-Effective	94.03	89.51	83.73	92.27
11	Codex (GPT-5.6 Terra Evidence)	95.60	86.69	77.36	92.26
12	OpenAI GPT-6 Astra	97.22	90.56	31.70	91.91
13	Codex (GPT-5.6 Luna Evidence)	93.50	87.43	80.05	91.16
14	Codex (GPT-6 Luna Evidence)	93.09	87.71	80.88	91.01
15	Reducto Deep Extract	94.20	80.47	92.01	90.44
16	Qwen3.8 Flash Next	94.82	87.81	37.74	89.88
17	Qwen3.8 27B	94.68	87.54	38.45	89.75
18	Extend (Max Context)	92.32	78.75	90.46	88.62
19	Qwen3.6 35B	92.86	86.37	36.90	88.11
20	Qwen3.5 35B	93.52	85.27	31.73	88.00
21	DeepSeek V4.1 Flash	94.44	81.49	22.23	87.11
22	Claude Code (Opus 4.8)	90.08	79.21	88.07	87.09
23	Claude Code (Opus 4.8 Evidence)	90.54	79.17	77.94	86.84
24	Qwen3.5 9B	92.31	81.12	32.22	86.10
25	Claude Code (Opus 5.5 Evidence)	89.63	77.96	80.83	86.06
26	Extend Extract	89.39	84.14	47.21	85.72
27	Datalab (Accurate + Balanced)	89.40	85.08	42.04	85.70
28	Reducto Extract	92.97	69.59	30.65	83.41
29	Kimi K3	94.64	69.64	4.99	83.17
30	Qwen3.5 4B	89.06	76.64	27.31	82.43
31	NuExtract3	86.41	77.01	32.71	81.02
32	GLM-5.3 Flash	96.30	51.56	27.83	80.75
33	Google Gemini 3.8 Flash	88.14	72.23	28.70	80.71
34	Google Gemini 3.5 Flash	87.87	69.76	27.90	79.84
35	Lift Datalab 9B	87.64	65.69	24.78	78.43
36	OpenAI GPT-5.4 Nano	77.43	76.37	35.81	74.90
37	Gemma4 26B	77.70	59.61	28.88	70.27
38	Gemma4 E4B	81.10	48.93	26.82	69.64
39	InternVL3.5 14B	80.70	39.04	19.41	66.35
40	InternVL3.5 8B	79.26	36.62	17.57	64.63
41	Qwen3.5 2B	71.22	50.93	25.93	63.40
42	Kimi-VL A3B	69.86	34.61	18.94	57.77
43	Gemma4 E2B	63.71	28.96	14.58	51.85
44	NuExtract 2.0 8B	62.92	23.50	6.39	49.42
45	Qwen3.5 0.8B	57.44	26.98	6.65	46.63
46	MiniCPM-V 4.5	57.39	25.30	1.13	45.85
47	Gemma4 12B	53.62	30.82	13.16	45.39
48	GLM-4.6V Flash	47.28	44.52	22.17	45.19

RESULTS

UNIKIE-BENCH

4,472 document images in 11 scenarios; macro F1 (%) by the benchmark's own evaluator. Published systems: UNIKIE-BENCH, Table 4. Equal scores share a rank.

#	System	Macro F1
1	**DocQuorum (ours)**	**82.65**
2	Gemini-3-Pro	82.37
3	Qwen3-VL-Plus	80.20
4	Qwen3-VL-8B	79.12
5	Qwen-VL-Max	77.77
6	GLM-4.1V-9B	69.91
7	GPT-4o	69.15
8	MiMo-VL-7B-RL	68.88
9	InternVL3.5-8B	67.70
10	GPT-5	67.59
11	MiniCPM-V4.5-8B	67.15
12	Claude-Sonnet-4.5	66.39
13	Kimi-VL-A3B	66.14
14	Gemma-3-12B	48.09
15	Ministral-3-8B	37.62
16	SmolVLM2-2.2B	23.69

RESULTS

VAREX

1,777 U.S. government forms, image input; exact match (%), normalized, by the benchmark's own scorer. Published systems: the VAREX-Bench image-modality leaderboard (20 models), which reports one decimal; ranks compare at that precision.

#	System	Score
1	**DocQuorum (ours)**	**98.23**
2	Gemini 2.5 Pro	98.0
3	Gemini 2.5 Flash	97.3
4	Qwen3-VL (8B)	96.6
5	Llama 4 Maverick (17B×128E)	95.6
6	Ministral (14B)	94.8
6	GPT-4o	94.8
8	Llama 4 Scout (17B×16E)	94.3
9	granite-vision-4.1-4b (4B)	94.2
10	NuExtract 2.0 (2B)	90.8
11	InternVL3.5 (2B)	85.6
12	granite-4.0-3b-vision (3B)	85.5
13	Qwen 2.5-VL (3B)	82.5
14	Gemma 3n (E2B)	71.0
15	MiniCPM-V4 (4B)	67.9
16	Gemma 3 (4B)	65.3
17	h2oVL Miss. (2B)	61.3
18	Qwen3-VL (2B)	34.2
18	h2oVL Miss. (800M)	34.2
20	InternVL3.5 (1B)	28.2
21	Qwen2-VL (2B)	9.7

RESULTS

UNIKIE-BENCH open category

1,661 documents in 8 scenarios (Chinese and English); macro F1 (%) by the benchmark's own evaluator. Published systems: the paper's open-category table.

#	System	Score
1	Gemini-3-Pro	81.65
2	**DocQuorum (ours)**	**72.33**
3	Qwen3-VL-Plus	70.61
4	Qwen-VL-Max	69.95
5	Qwen3-VL-8B	67.36
6	MiMo-VL-7B-RL	60.55
7	Kimi-VL-A3B	58.58
8	GLM-4.1V-9B	58.56
9	MiniCPM-V4.5-8B	52.60
10	InternVL3.5-8B	49.38
11	GPT-5	47.39
12	GPT-4o	46.69
13	Ministral-3-8B	41.58
14	Claude-Sonnet-4.5	39.71
15	Gemma-3-12B	30.90
16	SmolVLM2-2.2B	12.83

RESULTS

LongArray-Extract

45 long PDFs (bank statements, clinical listings, indictments); macro document accuracy (%) by the publisher's reference scorer. Published systems: Extend AI's results page, which reports one decimal; ranks compare at that precision.

#	System	Score
1	Extend MAX	99.2
2	**DocQuorum (ours)**	**98.05**
3	Reducto Deep Extract	97.4
4	Claude Opus 4.7	83.1
5	Reducto Standard	80.9
6	Pulse Effort	68.8
7	Pulse Auto	64.5
8	Gemini 3.1 Pro	47.3
9	LlamaParse Standard	47.2
10	LlamaParse Agentic	34.7
11	GPT-5.5	31.6
12	Gemini 3.5 Flash	31.1

RESULTS

DocuBench

72 documents in 10 file types; macro field accuracy (%) by the benchmark's own scorer. Published systems: the DocuBench leaderboard.

#	System	Score
1	DocuPipe (high effort)	97.02
2	DocuPipe (standard effort)	96.14
3	Claude Sonnet 5 (direct LLM)	91.73
4	**DocQuorum (ours)**	**91.54**
5	Reducto Deep Extract	89.38
6	Reducto (standard)	81.11
7	Extend	80.28
8	GPT-5.5 (direct LLM)	76.48
9	Gemini 3.5 Flash (direct LLM)	72.98
10	Unstructured	67.67